

Flexible Learning Reading Group @TU Berlin

6th Session: 6th of November 2019

SOLVING RUBIK'S CUBE WITH A ROBOT HAND

A PREPRINT

OpenAI

Ilge Akkaya*, Marcin Andrychowicz*, Maciek Chociej*, Mateusz Litwin*, Bob McGrew*, Arthur Petron*,
Alex Paino*, Matthias Plappert*, Glenn Powell*, Raphael Ribas*, Jonas Schneider*, Nikolas Tezak*,
Jerry Tworek*, Peter Welinder*, Lilian Weng*, Qiming Yuan*, Wojciech Zaremba*, Lei Zhang*

October 17, 2019




Picture credits - Akkaya et al. (2019: ArXiv)

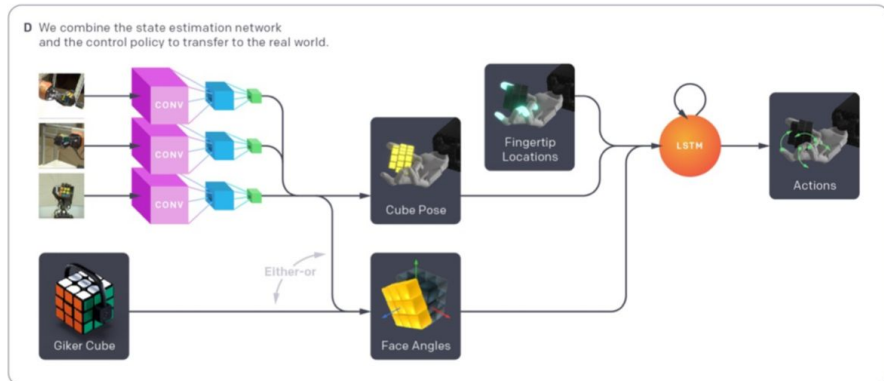


General Sim2Real Setup

Train in Simulation



Transfer to the Real World



A. Automatic Domain Randomization (ADR)

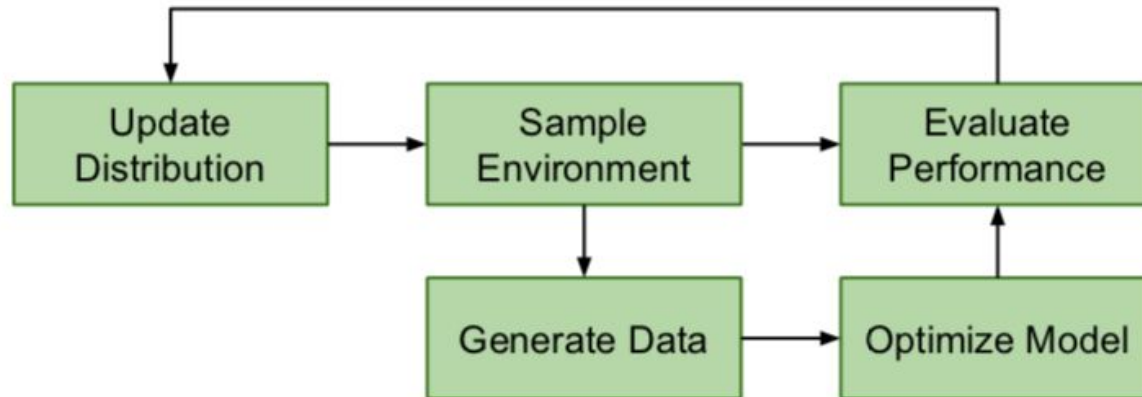
B. Distributed PPO + GAE LSTM-policy

C. ResNet-50-based State estimation

D. 0-Shot Sim2Real Transfer Results

Automatic Domain Randomization

Hypothesis: Transfer via Meta-Learning results from training on diverse set of environments



Limited network capacity enforces learn-to-learn!

➔ Natural Adaptive Curriculum!

On-Policy RL: Policy Gradient Methods

$$\pi^* = \arg \max_{\pi} V^{\pi}(s_0) = \arg \max_{\pi} \left[\mathbb{E}_{\pi} \left[\sum_t \gamma^{t-1} r_t | s_0 \right] \right]$$

→ Vanilla Policy Gradients (Sutton et al., 2000):

$$\nabla_{\theta} V^{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi} [\text{red X}(s, a) \nabla_{\theta} \ln \pi_{\theta}(a|s)]$$

→ Generalized Advantage Estimation (GAE):

$$\hat{V}_t^{(k)} = \sum_{i=t}^{t+k-1} \gamma^{i-t} r_i + \gamma^k V(s_{t+k})$$

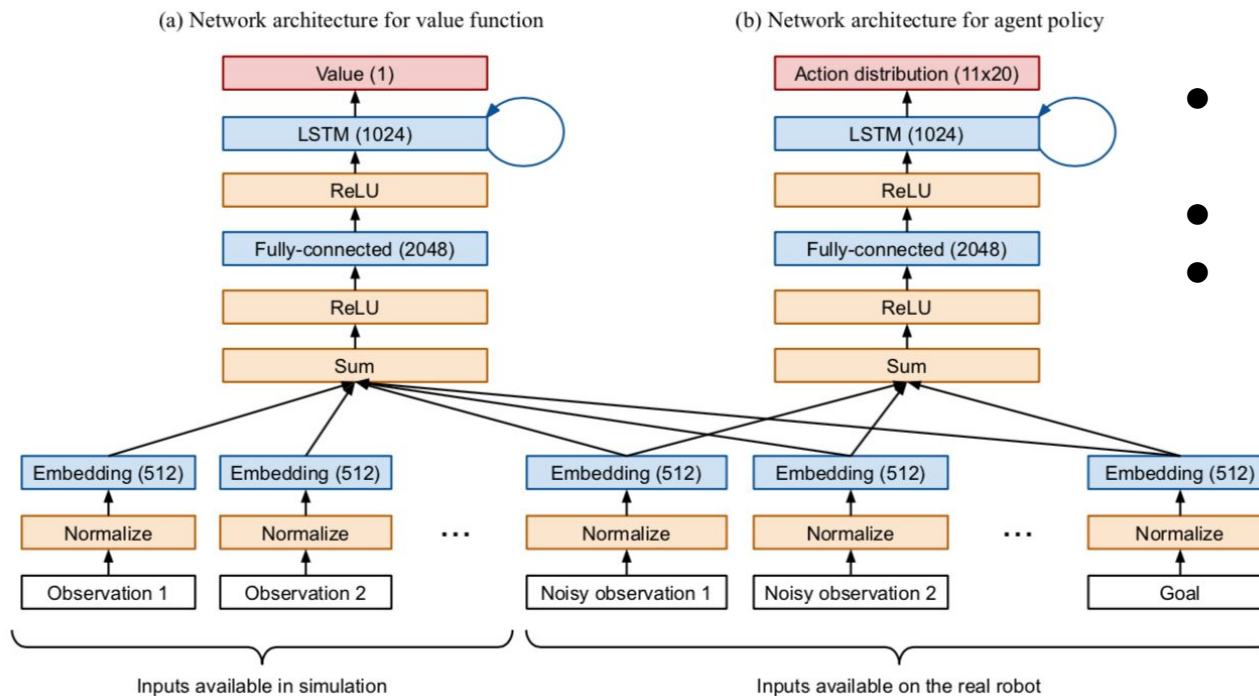
$$\hat{V}_t^{GAE} = (1 - \lambda) \sum_{k=0} \lambda^{k-1} \hat{V}_t^{(k)}, \quad 0 < \lambda < 1$$

\hat{A}_t^{GAE}

Training the Control Policy: PPO + GAE

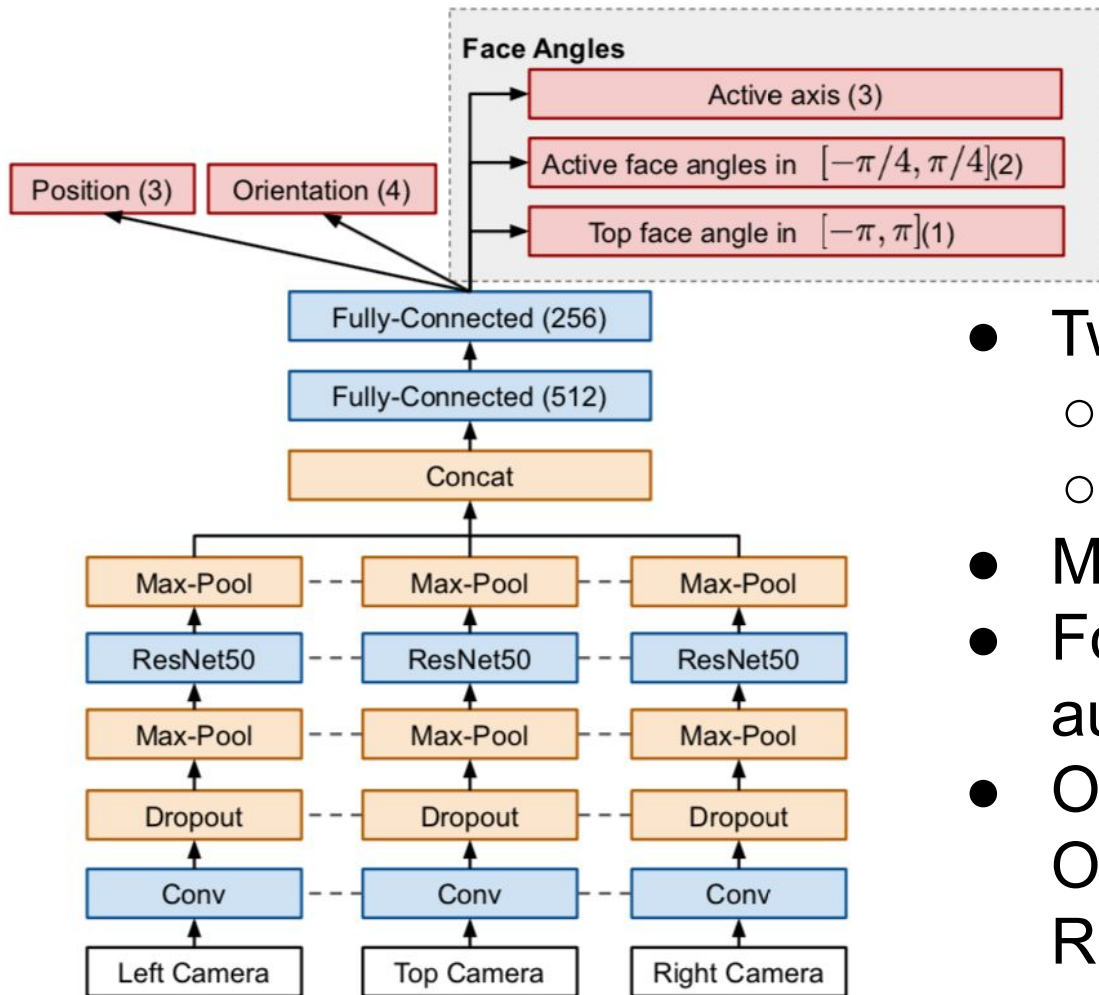
→ Proximal Policy Optimization (Schulman et al., 2017):

$$\max \mathbb{E} \left[\min \left(\frac{\pi(a_t|s_t)}{\pi_{old}(a_t|s_t)} \hat{A}_t^{GAE}, \text{clip} \left(\frac{\pi(a_t|s_t)}{\pi_{old}(a_t|s_t)}, 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t^{GAE} \right) \right]$$



- Value Net turned off during execution!
- 20 joints - 11 bins
- 'Multi-Categorical' Distribution

The State Estimation Vision System





- Two Versions:
 - Sticker Support
 - Sensors for Angles
- MSE + Cross-Entropy
- Focal Loss: Dynamic + automatic loss weights
- Only synthetic data: Open AI Remote Rendering Backend

Vision Module & Sim2Sim Results

Experiment	Errors (Sim)			Errors (Real)		
	Orientation	Position	Top Face	Orientation	Position	Top face
Full Model	6.52°	2.63 mm	11.95°	7.81°	6.47 mm	15.92°
No Domain Randomization	3.95°	2.97 mm	8.56°	128.83°	69.40 mm	85.33°
No Focal Loss	15.94°	5.02 mm	10.17°	19.10°	9.416 mm	17.54°
Non-discrete Angles	9.02°	3.78 mm	42.46°	10.40°	7.97 mm	35.27°


➔ ADR + Adaptive Weighting of different tasks is crucial




➔ Curriculum outperforms fixed levels of randomness

Sim2Real Results


Policy	Sensing			Successes (Real)		Success Rate	
	Pose	Face Angles	ADR Entropy	Mean	Median	Half	Full
Manual DR	Vision	Giiker	-0.569^* npd	1.8 ± 0.4	2.0	0 %	0 %
ADR	Vision	Giiker	-0.084 npd	3.8 ± 1.0	3.0	0 %	0 %
ADR (XL)	Vision	Giiker	0.467 npd	17.8 ± 4.2	12.5	30 %	10 %
ADR (XXL)	Vision	Giiker	0.479 npd	26.8 ± 4.9	22.0	60 %	20 %
ADR (XXL)	Vision	Vision	0.479 npd	12.8 ± 3.4	10.5	20 %	0 %




(a) Unperturbed (for reference).




(b) Rubber glove.




(c) Tied fingers.



(d) Blanket occlusion and perturbation.



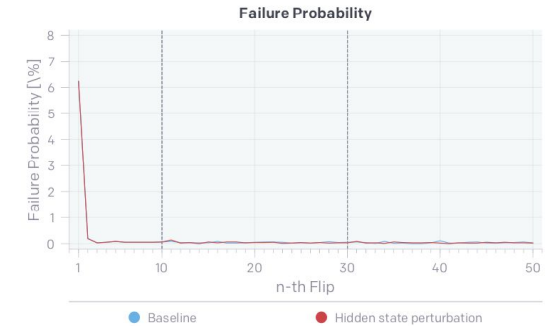
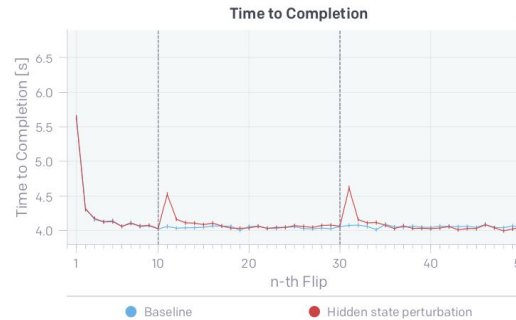


(e) Plush giraffe perturbation.¹⁷



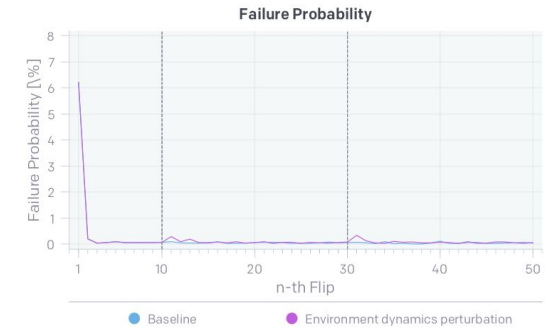
(f) Pen perturbation.

LSTM + ADR = Meta-Learning?

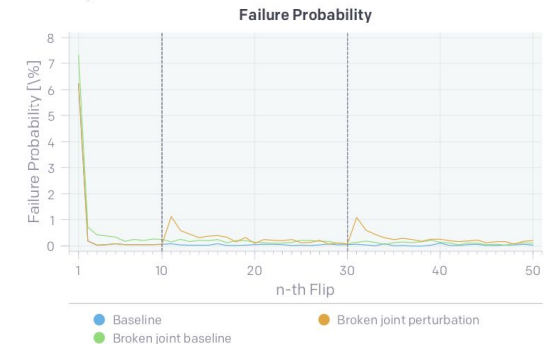
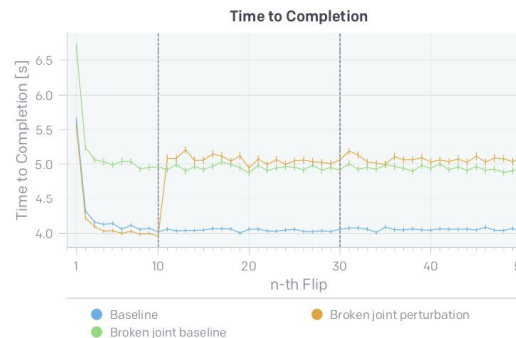
Training



(a) Resetting the hidden state.



(b) Re-sampling environment dynamics.



(c) Breaking a random joint.

Cool things that I learned!

1. ORRB remote rendering backend on top of Unity
2. 'Bi-directional' ADR: Entropy \downarrow if performance $<$ thresh
3. Redis: Centralized storage of parameters + data
4. Adversarial Random Networks - Exploration
5. Concat + Add Embedding - Flexible Inputs
6. Policy Distillation to transfer progress
7. Face Angle "Classification" - Cross-Entropy Loss
8. Multi-Task Vision with Focal Loss Weighting
9. LARS optimizer - Large Batchsize 1024



Gary Marcus
@GaryMarcus

Since @OpenAI still has not changed misleading blog post about "solving the Rubik's cube", I attach detailed analysis, comparing what they say and imp they actually did. IMHO most would not be nonexperts.

Please zoom in to read & judge for yourself

Reality

- Neural networks didn't do the solving; a 17-year old symbolic AI algorithm did
- The solving (which face should turn where) algorithm was innate, not learned.
- Reinforcement learning played no role in the choice of which faces to turn (ie what most people call solving).
- What was learned was object manipulation, not cube solving
- Only ONE object was manipulated, and there was no test of generalizability to other objects
- That object was heavily instrumented (eg with bluetooth sensors). The hand was instrumented with LEDs, as well.
- Success rate was only 20%; hand frequently dropped cube